Few-shot label effects: retain matched zero-support controls

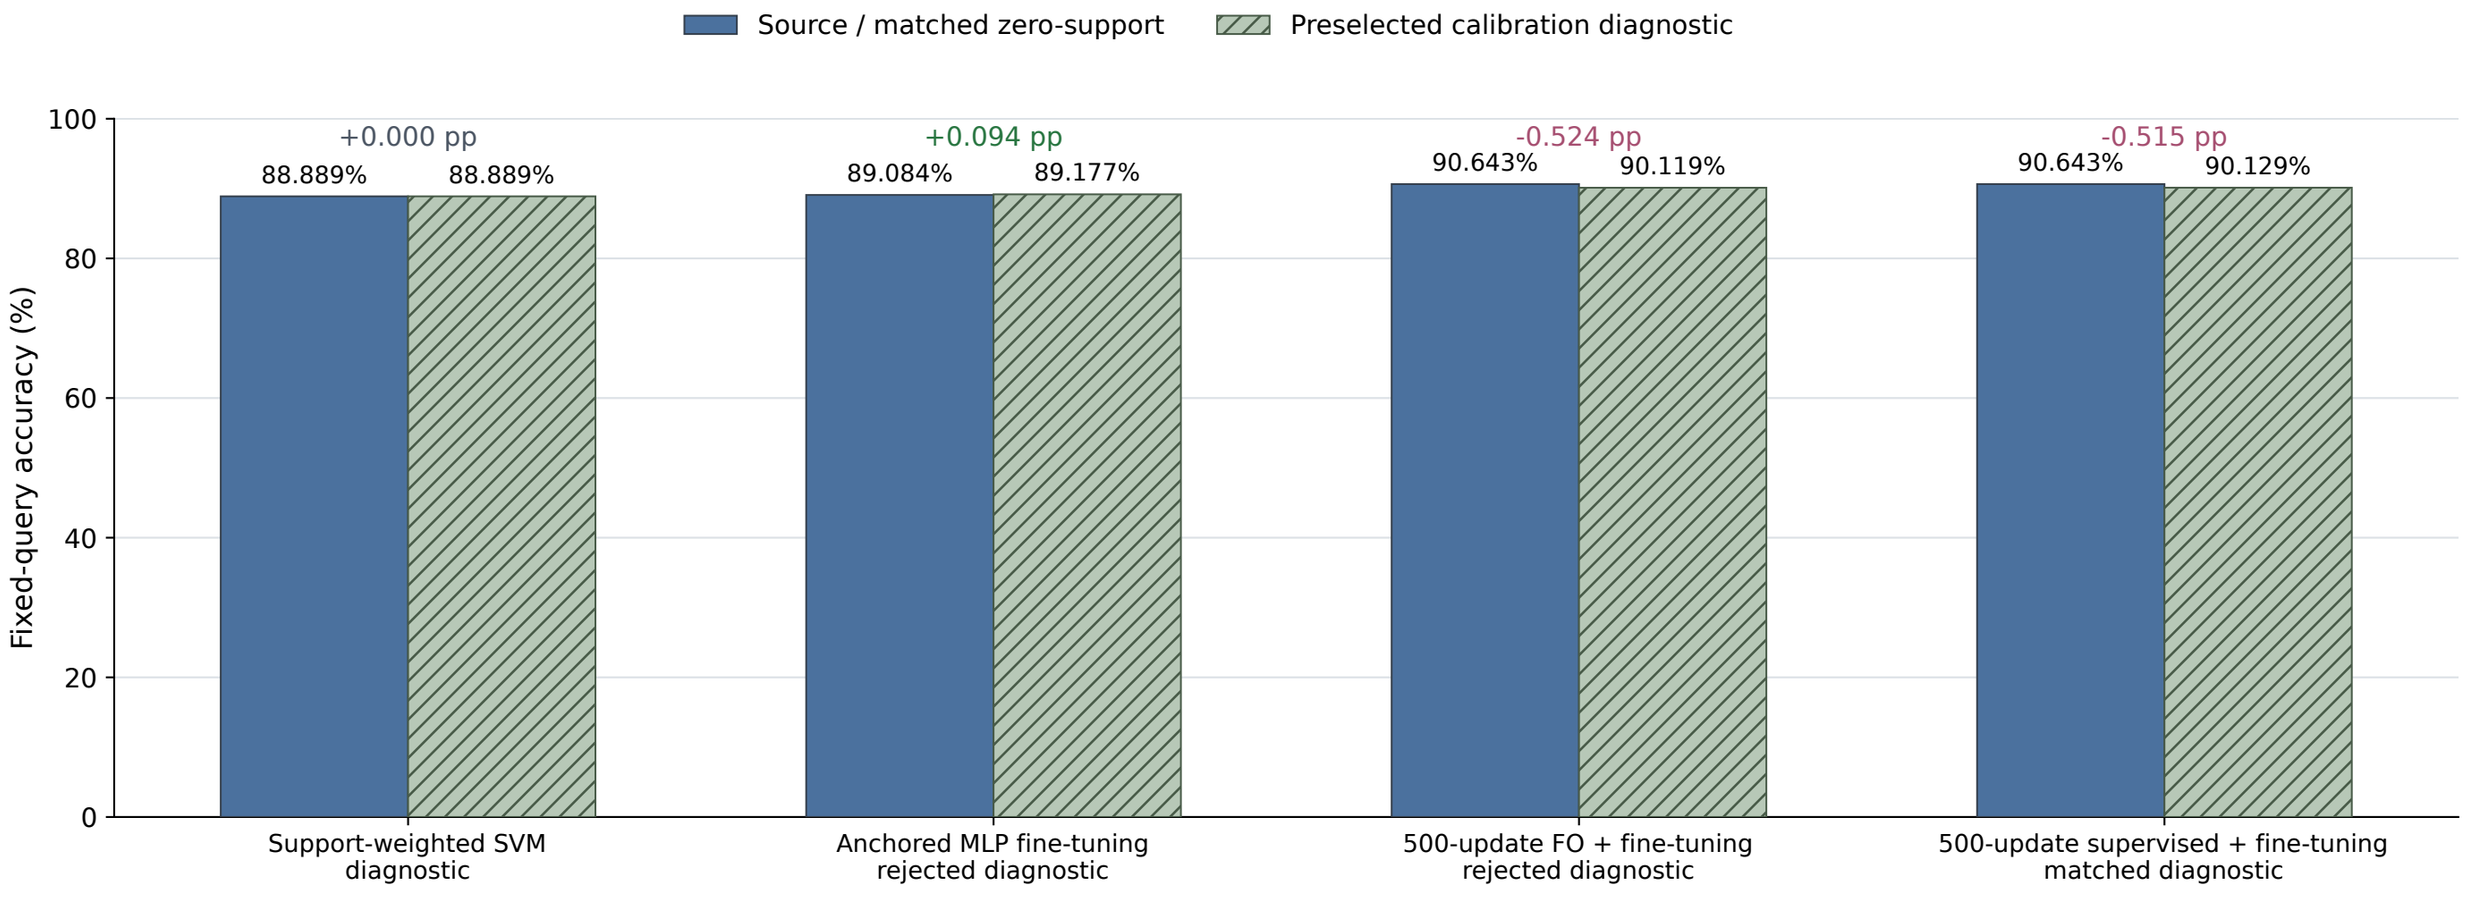

Each labelled-support draw uses 1 independent specimen per anonymous class (7 labels total); the query stays fixed at 171 specimens.

No support-based calibration passed the full development gate in R2 / R3 / R4: all retained their zero-support source predictor.

The 500-update meta rows are diagnostics. R4 retained the 100-update source, with 90.2534% accuracy and no target calibration.

Neural values average 3 seeds × 100 support draws; SVM uses 100 draws. These repeated cases are not independent samples or batches.

Single-batch retrospective results on historically explored data. Axes start at zero; no statistical confidence intervals.