# Source-model improvements and supervised controls

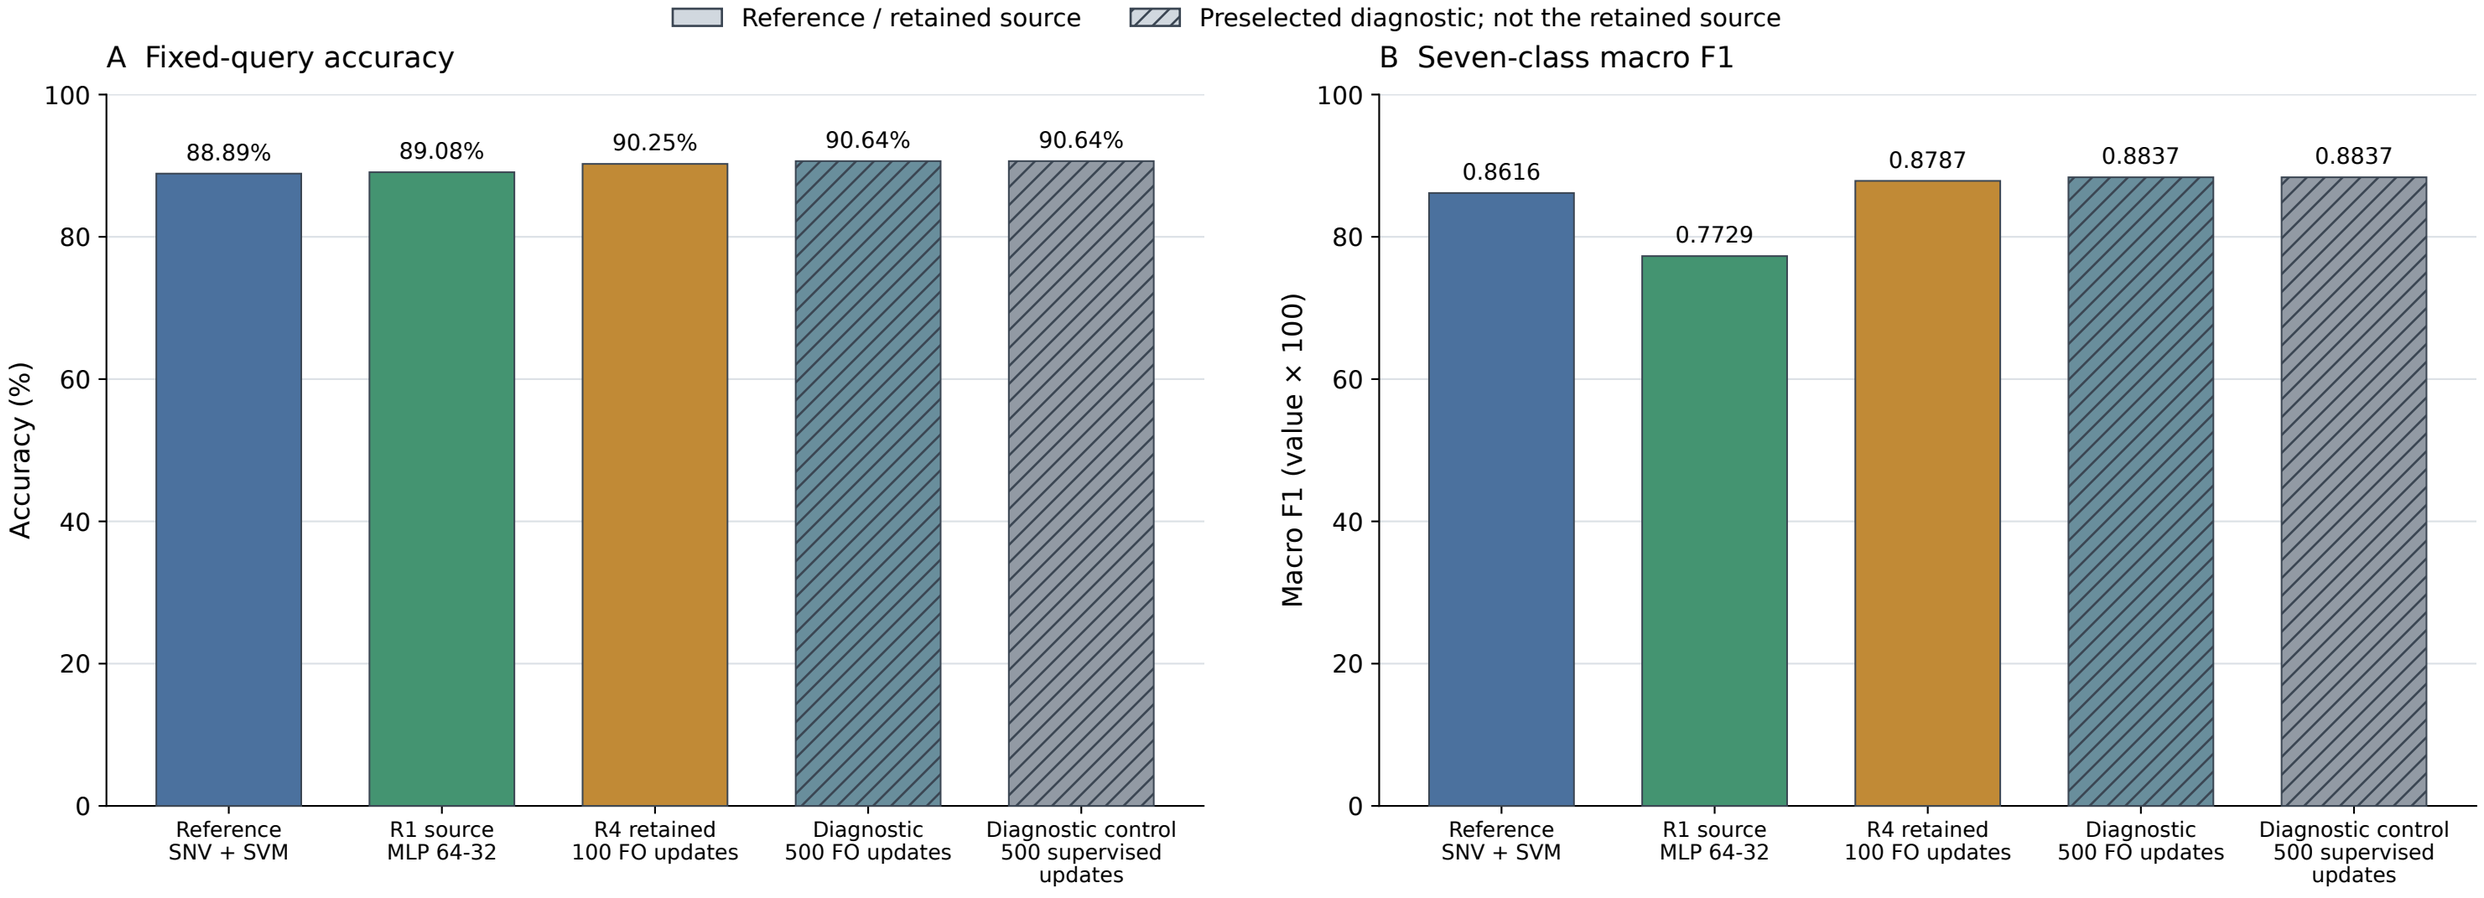


All source models use 513 training specimens and zero target-calibration labels; fixed query n = 171.

Neural bars average seeds 42 / 43 / 44; SVM is the fixed reference. Diagnostic status was locked using development data.

FO and matched supervised source models make identical query classifications at both budgets for all 3 seeds; logits differ.

Single-batch retrospective study; the same historical query specimens are reused. Source gains are not calibration gains.

Axes start at zero. No statistical confidence intervals; random seeds are not independent acquisition batches.