Paper-inspired classification pilot: selected and diagnostic outcomes

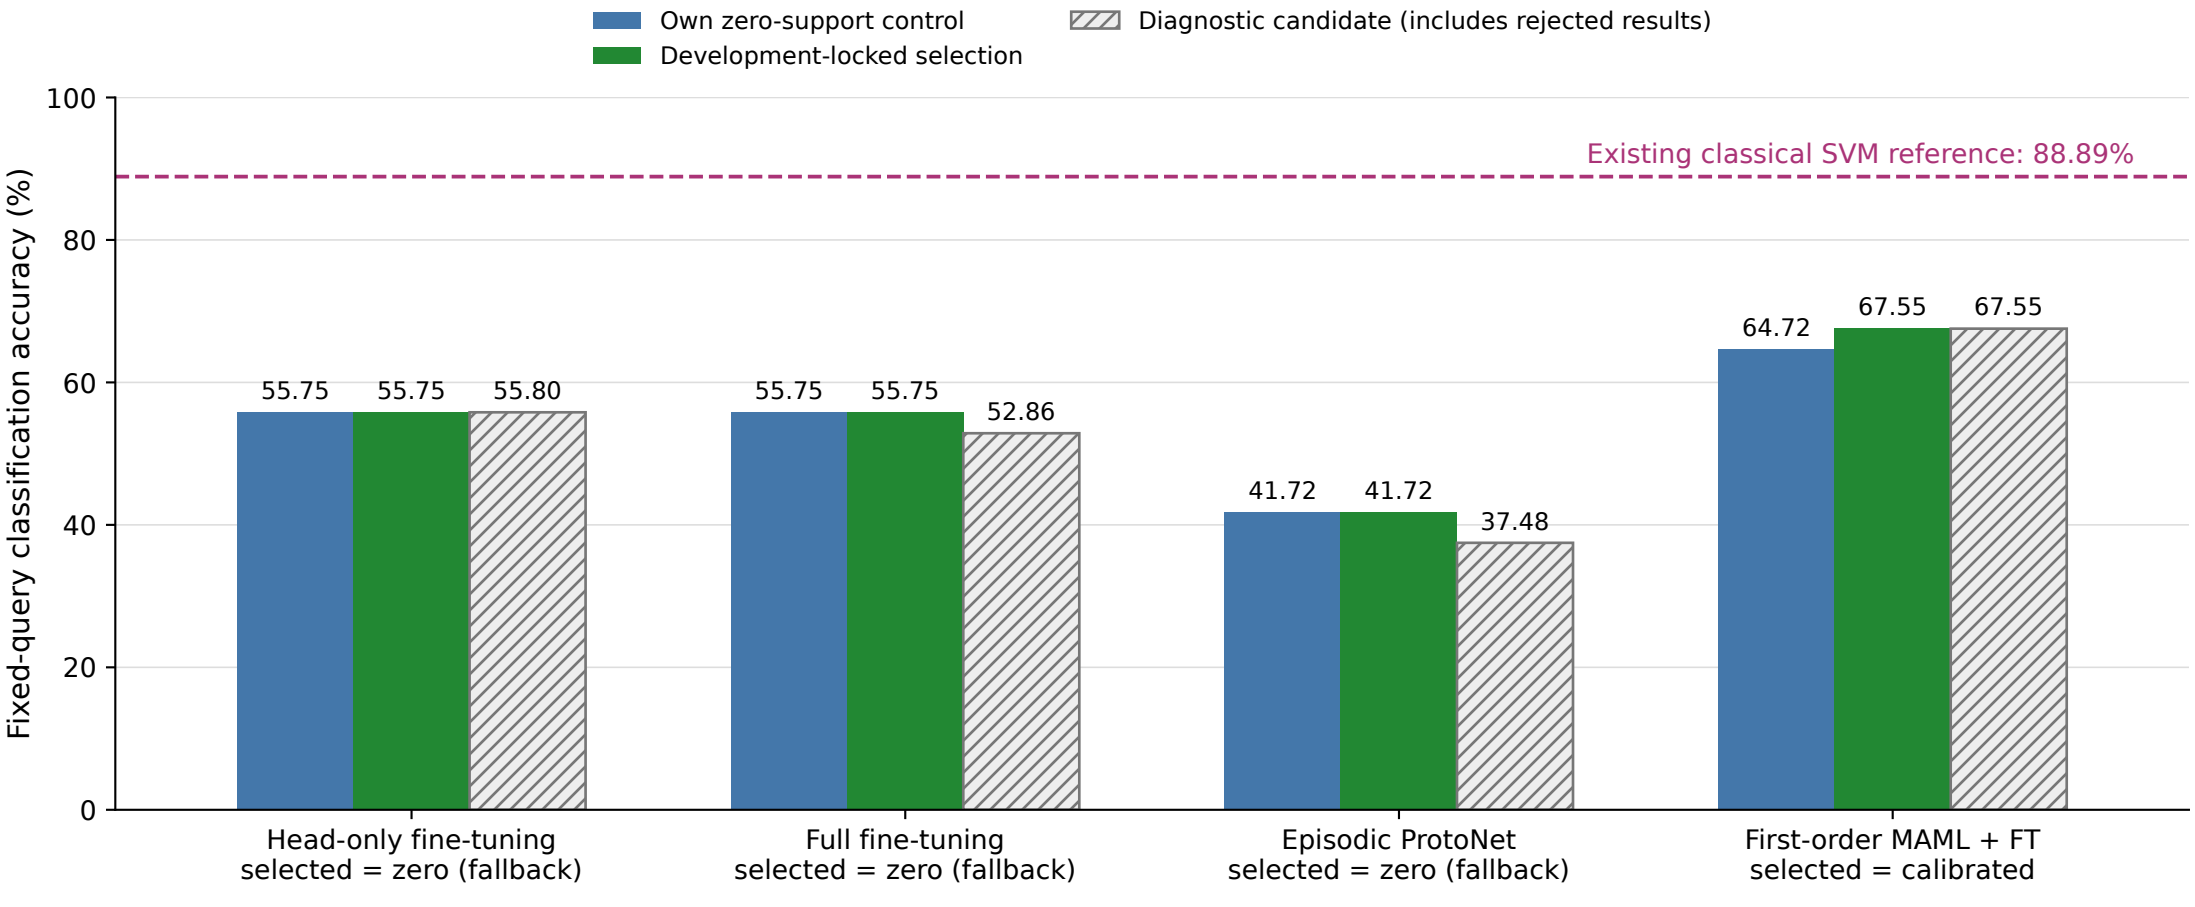

Bars: means over 3 seeds × 100 support draws; each draw labels 1 specimen per class (7 total).

Each strategy uses its own zero-support control; differences between source models are not calibration gains.

Single-batch retrospective evaluation; the same 171 query specimens are reused in all 300 non-independent cases.

No statistical confidence intervals. Paper-inspired classification adaptation; the original benchmark studied regression.